

Deep Learning

An Introduction

Compiled by

Dr. Sandeep Palakkal

Samsung R&D Institute India, Bangalore

Dec 14, 2018

SAMSUNG

Machine Learning

SAMSUNG

© Samsung 2018

“Machine learning gives computers the ability to learn without being explicitly programmed.” -- Arthur Samuel, 1959.

[Compose Mail](#)

[Inbox \(5662\)](#)

[Sent Mail](#)

[All Mail](#)


[Spam \(316\)](#)



<https://en.wikipedia.org>



<http://www.image-net.org>



<https://en.wikipedia.org>



<https://www.trademed.com>



<https://www.expertsystem.com>



<https://www.techleer.com>



<https://services.harman.com/>

Customers who bought this item also bought



Reality Is Not What It Seems: The Journey to Quantum Gravity
• Carlo Rovelli
★★★★☆ 12
Kindle Edition
₹ 360.39



The Order of Time
• Carlo Rovelli
★★★★★ 17
Kindle Edition
₹ 517.18



Six Easy Pieces: Essentials of Physics Explained by Its Most Brilliant Teacher
• Richard P. Feynman
★★★★★ 28
Kindle Edition
₹ 788.22



What is time? What is space? (I Dialoghi)
• Carlo Rovelli
Kindle Edition
₹ 471.92




Origin Story: A Big History of Everything
• David Christian
★★★★★ 18
Kindle Edition
₹ 350.46

<https://www.amazon.in>




<https://towardsdatascience.com>





<https://www.samsung.com>

Machine Learning : The Conventional Approach

Conventional ML



Example Human Detection using HOG [1]




❖ Designing SIFT took 10 years of effort by Prof. David G. Lowe [2]

❖ Can we do better?

[1] N. Dalal and B. Triggs, "Histograms of Oriented Gradients for Human Detection," *CVPR* 2005.

[2] D. G. Lowe, "Distinctive image features from scale-invariant keypoints," *IJCV* 2004

Deep Learning : What?



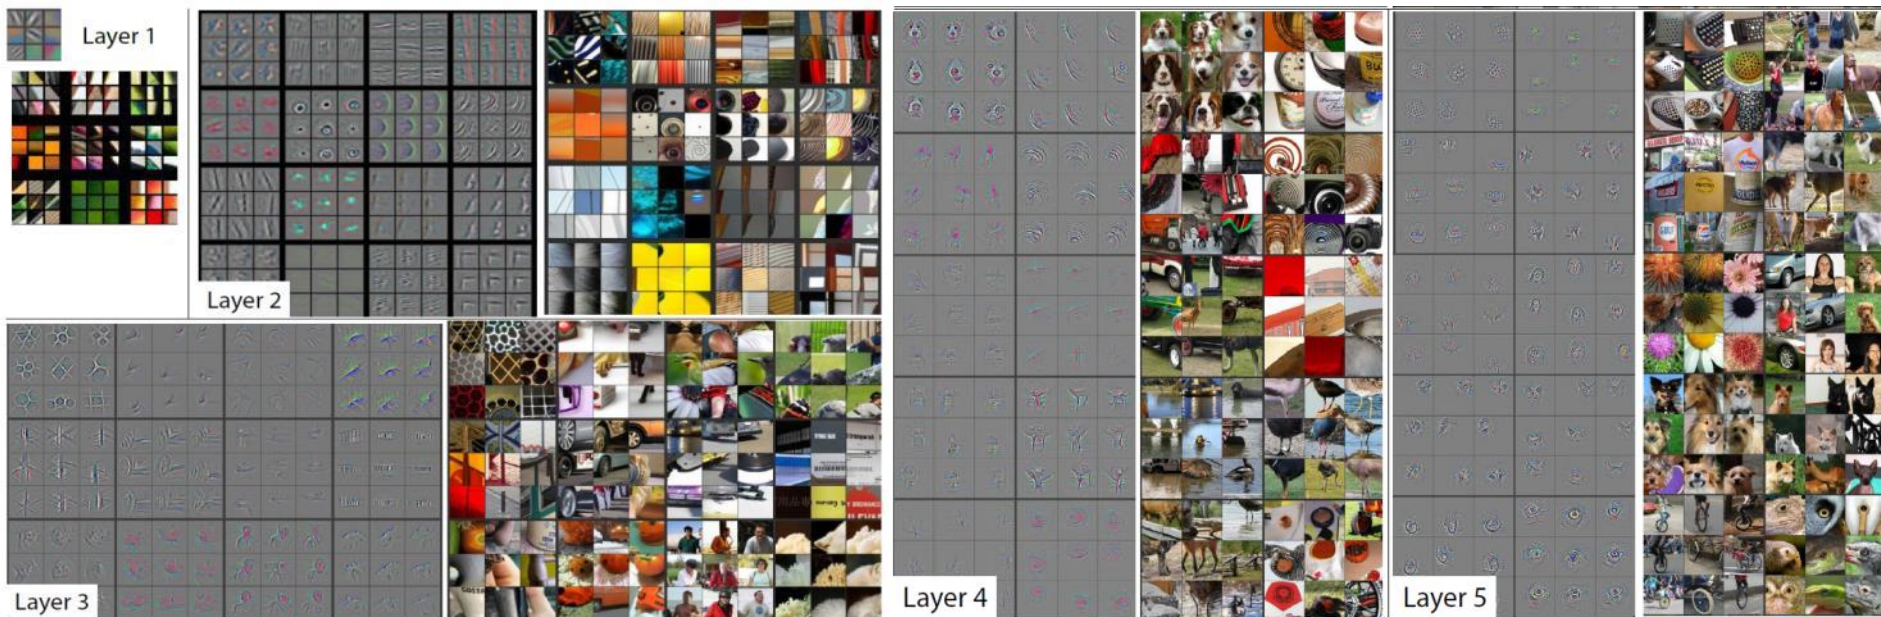
Several Areas of Influence

- ✓ Affective Computing
- ✓ Artificial Intelligence
- ✓ Autonomous Cars
- ✓ Biology
- ✓ Business
- ✓ Computer Vision
- ✓ Finance
- ✓ Gaming
- ✓ Genomics
- ✓ Healthcare/Health Informatics
- ✓ Information Retrieval
- ✓ Natural Language Processing
- ✓ Sparse Coding
- ✓ Speech Processing
- ✓ Trading
- ✓ Weather Forecast

- Deep learning is a **machine learning** approach that makes use of **deep neural networks (DNN)**
- A DNN consists of a hierarchy of computational layers (usually, more than 2 layers)
- Each layer in DNN transforms the input data into slightly higher and more abstract representation
- The idea of hierarchical representation and DNN architectures are inspired by Visual Cortex of brain

Features Learned by AlexNet for Image Classification

- DNN learns features in multiple layers of abstraction
- Level of abstraction increases as we go deeper
- Deeper layers capture more complex structures pertinent to data



Deep V/S Shallow Learning



Courtesy: <http://cs231n.github.io/classification/>



❑ Features: **Selectivity vs invariance**

- *Selective* to relevant variations (e.g., between class)
- *Invariant* to irrelevant variations (e.g., within-class)

❑ Deep networks: complex, hierarchical learning


- Stack of non-linear input-output mapping
- Each stack transforms input to increase selectivity & invariance

❑ Shallow networks would be heavier to get same performance




Artificial Neural Networks (ANN)

Single Perceptron (Rosenblatt-1958)




Activation function	Equation	Example	1D Graph
Unit step (Heaviside)	$\phi(z) = \begin{cases} 0, & z < 0, \\ 0.5, & z = 0, \\ 1, & z > 0, \end{cases}$	Perceptron variant	
Sign (Signum)	$\phi(z) = \begin{cases} -1, & z < 0, \\ 0, & z = 0, \\ 1, & z > 0, \end{cases}$	Perceptron variant	
Linear	$\phi(z) = z$	Adaline, linear regression	
Piece-wise linear	$\phi(z) = \begin{cases} 1, & z \geq \frac{1}{2}, \\ z + \frac{1}{2}, & -\frac{1}{2} < z < \frac{1}{2}, \\ 0, & z \leq -\frac{1}{2}, \end{cases}$	Support vector machine	
Logistic (sigmoid)	$\phi(z) = \frac{1}{1 + e^{-z}}$	Logistic regression, Multi-layer NN	
Hyperbolic tangent	$\phi(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$	Multi-layer NN	

Logistic Regression/Classification by Perceptron



The Neural Network: Multilayer Perceptrons (MLP)



A 2-Layer MLP (Single Hidden Layer) is a universal approximator!!!

Classification Boundary of a 2-Layer MLP







Image Credit: <http://cs231n.github.io/neural-networks-1/>

Training Neural Networks: Gradient Descent & Backpropagation

Objective of training: Find weights w that reduces an error/loss between output and desired output



Loss Function

❑ Desired output: $\mathbf{y} = \{y_1, y_2, \dots, y_m\}$

❑ Output of network: $\hat{\mathbf{y}} = \{\hat{y}_1, \hat{y}_2, \dots, \hat{y}_m\}$

❑ Loss functions:

▪ Euclidean: $L = \frac{1}{2m} \sum_{j=1}^m (\hat{y}_j - y_j)^2$

❑ Other loss functions:

▪ Entropy loss function, L1 Loss function, etc.

Gradient Descent & backpropagation

• Iterative algorithm


• k-th iteration: $w^k = w^{k-1} - \eta \frac{\partial L}{\partial w}$

• Apply chain rule to compute gradient:

$$\frac{\partial L}{\partial w_i} = \frac{\partial L}{\partial y_j} \frac{\partial y_j}{\partial u_j} \frac{\partial u_j}{\partial w_i}$$

• For hidden layers, apply chain rule backwards

• Gradient of later layers propagates backward [1]



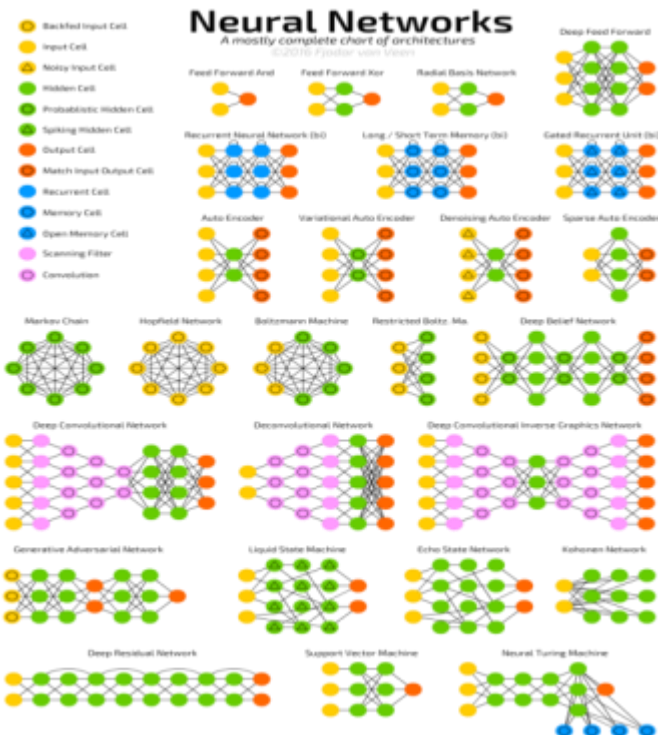
[1] Tom M. Mitchell, Machine Learning, McGraw Hill, 1997

[2] A. Karpathy, Stanford University CS231N Course Notes, <http://cs231n.github.io/>

Types of DNNs

Several types of (deep) neural networks

- Convolutional neural networks (CNN)
- Recurrent neural networks (RNN)
- Deep belief networks (DBN)
- Auto-encoders
- Long short-term memory (LSTM)
- (Restricted) Deep Boltzmann networks (DBM & RBM)



Convolutional Neural Networks (CNN)




Image credit: M. Peemen et al., VLIW Code Generation for a Convolutional Network Accelerator, SCOPES, 2015

RNN & LSTM




Image credits: Y. LeCun, Y. Bengio and G. Hinton, Deep Learning, *Nature* 2015, <http://deeplearning.net/tutorial/lstm.html>

Autoencoder





Image credits: Wikipedia

Convolutional Neural Networks (CNN)


- 3 Key ideas :
 - Local receptive field
 - Shared weights
 - Spatial/temporal subsampling (pooling)



Convolution in 3D (Image courtesy: <https://ireneli.eu/>)

FULLY CONNECTED NEURAL NET


Example: 1000x1000 image
1M hidden units
➡ 10¹² parameters!!!



- Spatial correlation is local
- Better to put resources elsewhere!

LOCALLY CONNECTED NEURAL NET


Example: 1000x1000 image
1M hidden units
Filter size: 10x10
100M parameters




Ranzai

Convolutional Neural Networks (CNN)


- 3 Key ideas :
 - Local receptive field
 - Shared weights
 - Spatial/temporal subsampling (pooling)



Convolution in 3D (Image courtesy: <https://ireneli.eu/>)



3 Key ideas in a Convolution Layer (Image Courtesy: Cloudy Ngyuen)



CNN Architecture (Image courtesy: <https://ireneli.eu/>)

A Primer on Convolution

Input image





Image credit: Wikipedia

Convolution Kernel

$$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$$

Feature map




1 _{x1}	1 _{x0}	1 _{x1}	0	0
0 _{x0}	1 _{x1}	1 _{x0}	1	0
0 _{x1}	0 _{x0}	1 _{x1}	1	1
0	0	1	1	0
0	1	1	0	0

Image

Source: <http://deeplearning.stanford.edu/>

4		

Convolved Feature



Convolution in 3D



Source: B. P. Lathi, Linear Systems and Signals, 2nd Ed., 2004


- The most fundamental operation in signal/image processing
- For an $M \times N$ image

$$y(m,n) = \sum_{p=0}^{M-1} \sum_{q=0}^{N-1} x(p,q)h(m-p,n-q)$$

- Measures “similarity” between image and convolutional filter
- In CNN (3D):

$$y(m,n,k) = \sum_{c=0}^{C-1} \sum_{p=0}^{M-1} \sum_{q=0}^{N-1} x(p,q,c)h(m-p,n-q,c)$$

Evolution of CNN : Summary




Sparse connections : S2 to C3

	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
0	X				X	X	X			X	X	X	X	X	X	X
1	X	X	X		X	X	X			X	X	X	X	X	X	X
2	X	X	X			X	X	X		X	X	X	X	X	X	X
3		X	X	X		X	X	X		X	X	X	X	X	X	X
4			X	X	X		X	X	X	X	X	X	X	X	X	X
5				X	X	X		X	X	X	X	X	X	X	X	X


Le-Net

Le-Cun
1998




Krizhevsky
2012

AlexNet Architecture




Performance on
ImageNet
Classification

A. Canziani et. al, "An Analysis of Deep Neural Network Models for Practical Applications." *arXiv preprint* 2016




Inception Module (GoogLeNet)

zegey 2014




VGG Net

Simonyan
2014



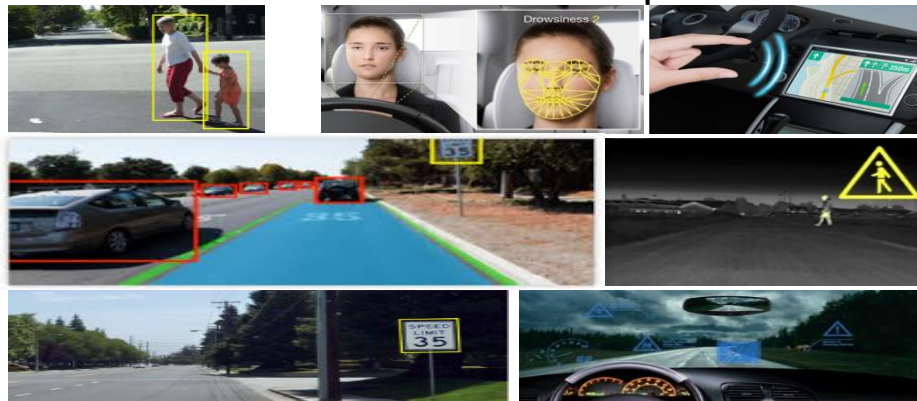
Residual learning
(ResNet)

He 2015




Deep Learning for Autonomous Cars

Autonomous Cars Use Case Scenario Examples



- Pedestrian/cyclist detection
- Free space and lane detection
- Vehicle detection
- Driver monitoring
- Traffic sign board recognition
- Differentiation of vehicle types
- HD mapping of road, etc.

Deep Learning Challenges in Autonomous Cars



- Highly accurate ML models
- State-of-the-art performance in ADAS



- DNNs are large with very high computational complexity requirements
- Large memory requirement on embedded systems
- Multiple optimizations required for deploying on embedded system

